

Recipe Recommendation assignment :-

By Gaurav Chaudhary, Gaganpreet Kaur, Malaika Goveas (DSC- 59)

Problem Statement

Step into the shoes of an ML engineer working at food.com. Your job is to design a recommender system to recommend recipes to users based on their choice and the current recipe they are looking at.

The recommendation engine is a way to increase the website's user engagement.

If a user is shown relevant recipes, they are more likely to spend more time on your site reading about recipes.

Higher user engagement will likely result in more business opportunities like collaborations, promotions, etc.

The performance of a recommendation engine will significantly impact the revenue your recipe site can generate.

Designing a recommender from scratch is a time-consuming task.

In this assignment, you are expected to explore the data and create features that will be used to build the recommender.

You will be working with the two CSV files linked below.

1. **Raw_recipes_cleaned.csv** - The first file is the Raw_recipes.csv file. It contains all the recipe-related information. Each row in this file describes a recipe
2. **RAW_interactions_cleaned.csv** - The second file we will be using is the RAW_interactions.csv. Each row in this data file is one user reviewing one recipe. One user can review more than one recipe, and each recipe can be reviewed by more than one user, so there is a many-to-many relationship between users and recipes, but the combination of user_id and reviewer_id in each row will be unique.

Task List

Task 1: Read the data

1. Read RAW_recipes.csv from S3 bucket.
2. Ensure each field has the correct data type.

name	id	minutes	contributor_id	submitted	tags	nutrition	n_steps	steps	description	ingredients	n_ingredients
arriba baked wi...	137739	55	47892	2005-09-16 00:00:00	['60-minutes-or-l...	[51.5, 0.0, 13.0, ...]	11	['make a choice a...	autumn is my favo...	['winter squash', ...]	7
a bit different ...	31490	30	26278	2002-06-17 00:00:00	['30-minutes-or-l...	[173.4, 18.0, 0.0...]	9	['preheat oven to...	this recipe calls...	['prepared pizza ...]	6
all in the kitche...	112140	130	196586	2005-02-25 00:00:00	['time-to-make', ...]	[269.8, 22.0, 32....]	6	['brown ground be...	this modified ver...	['ground beef', '...]	13
alouette potatoes	59389	45	68585	2003-04-14 00:00:00	['60-minutes-or-l...	[368.1, 17.0, 10....]	11	['place potatoes ...]	this is a super e...	['spreadable chee...	11
amish tomato ket...	44061	190	41706	2002-10-25 00:00:00	['weeknight', 'ti...	[352.9, 1.0, 337....]	5	['mix all ingredi...	my dh's amish mot...	['tomato juice', ...]	8

Task 2: Extract individual features from the nutrition column.

id	nutrition	calories	total fat (PDV)	sugar (PDV)	sodium (PDV)	protein (PDV)	saturated fat (PDV)	carbohydrates (PDV)
137739	51.5, 0.0, 13.0, 0.0, 2.0, 0.0, 4.0	51.5	0.0	13.0	0.0	2.0	0.0	4.0
31490	173.4, 18.0, 0.0, 17.0, 22.0, 35.0, 1.0	173.4	18.0	0.0	17.0	22.0	35.0	1.0
112140	269.8, 22.0, 32.0, 48.0, 39.0, 27.0, 5.0	269.8	22.0	32.0	48.0	39.0	27.0	5.0
59389	368.1, 17.0, 10.0, 2.0, 14.0, 8.0, 20.0	368.1	17.0	10.0	2.0	14.0	8.0	20.0
44061	352.9, 1.0, 337.0, 23.0, 3.0, 0.0, 28.0	352.9	1.0	337.0	23.0	3.0	0.0	28.0

Task 3: Standardize the nutrition values.

Convert the nutritional values to per 100 calories.

id	total fat (PDV)	sugar (PDV)	sodium (PDV)	protein (PDV)	saturated fat (PDV)	carbohydrates (PDV)
137739	0.0	13.0	0.0	2.0	0.0	4.0
31490	18.0	0.0	17.0	22.0	35.0	1.0
112140	22.0	32.0	48.0	39.0	27.0	5.0
59389	17.0	10.0	2.0	14.0	8.0	20.0
44061	1.0	337.0	23.0	3.0	0.0	28.0

only showing top 5 rows

id	total_fat_per_100_cal	sugar_per_100_cal	sodium_per_100_cal	protein_per_100_cal	saturated_fat_per_100_cal	carbohydrates_per_100_cal
137739	0.0	25.24271844660194	0.0	3.883495145631068	0.0	7.766990291262136
31490	10.380623202758338	0.0	9.80392191371621	12.687428358926859	20.18454511647455	0.5767012890421299
112140	8.154188656554616	11.860638045897625	17.79095706884644	14.455152618437731	10.007413351226122	1.8532246946715039
59389	4.618310165205302	2.71665303835606	0.543330607671212	3.8033142536984843	2.173322430684848	5.43330607671212
44061	0.2833663976467306	95.49447600694822	6.517427145874804	0.8500991929401919	0.0	7.934259134108458

Task 4: Convert the tags column from a string to an array of strings.

id	tags
137739	[60-minutes-or-less, time-to-make, course, main-ingredient, cuisine, preparation, occasion, north-american, side-dishes, vegetables, mexican, easy, fall, holiday-event, vegetarian, winter, dietary, chris
31490	[30-minutes-or-less, time-to-make, course, main-ingredient, cuisine, preparation, occasion, north-american, breakfast, main-dish, pork, american, oven, easy, kid-friendly, pizza, dietary, northeastern-un
112140	[time-to-make, course, preparation, main-dish, chili, crock-pot-slow-cooker, dietary, equipment, 4-hours-or-less]
59389	[60-minutes-or-less, time-to-make, course, main-ingredient, preparation, occasion, side-dishes, eggs-dairy, potatoes, vegetables, oven, easy, dinner-party, holiday-event, easter, cheese, stove-top, dieta
44061	[weeknight, time-to-make, course, main-ingredient, cuisine, preparation, occasion, north-american, canning, condiments-etc, vegetables, american, heirloom-historical, holiday-event, vegetarian, dietary,

only showing top 5 rows

tags column is a ArrayType(StringType(), False)

Task 5: Read the second data file

Read the RAW_interaction.csv and join this interaction level file with the recipe level data frame.



```
(interaction_level_df.count() ,len(interaction_level_df.columns) )
```

```
(1132367, 30)
```

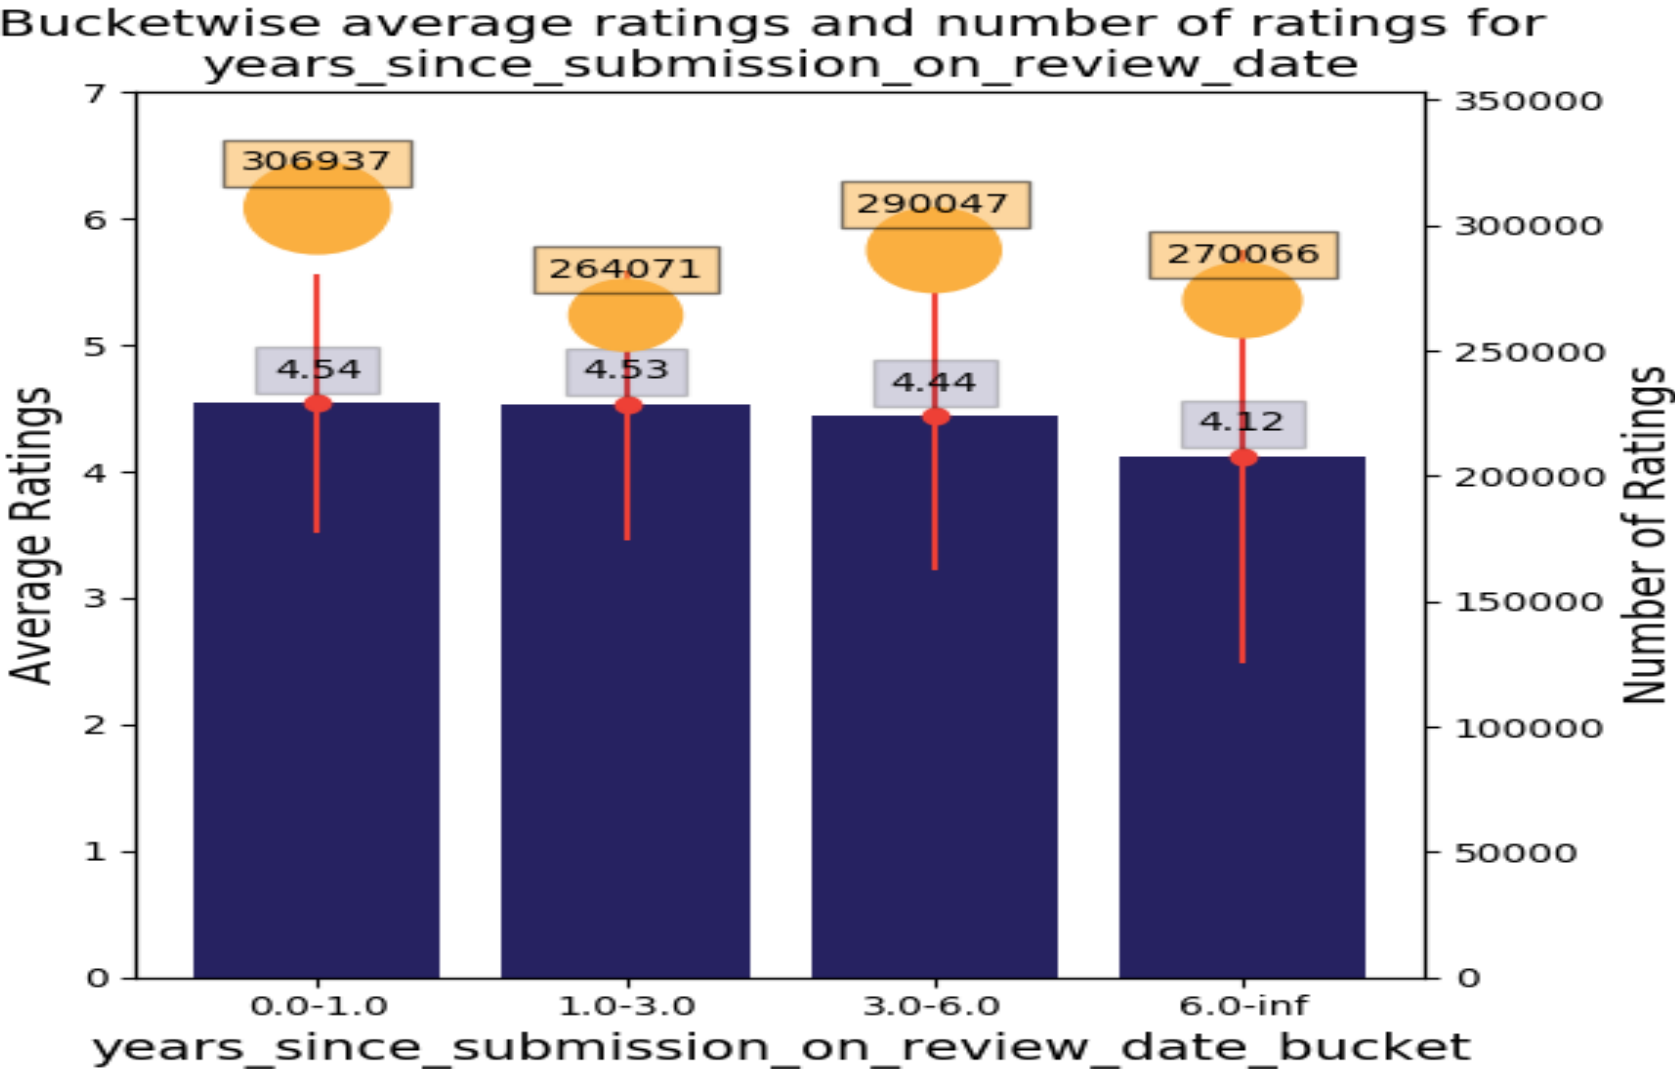
Task 6: Create time-based features.

Create features that capture the time passed between one review and the date on which the recipe was submitted.

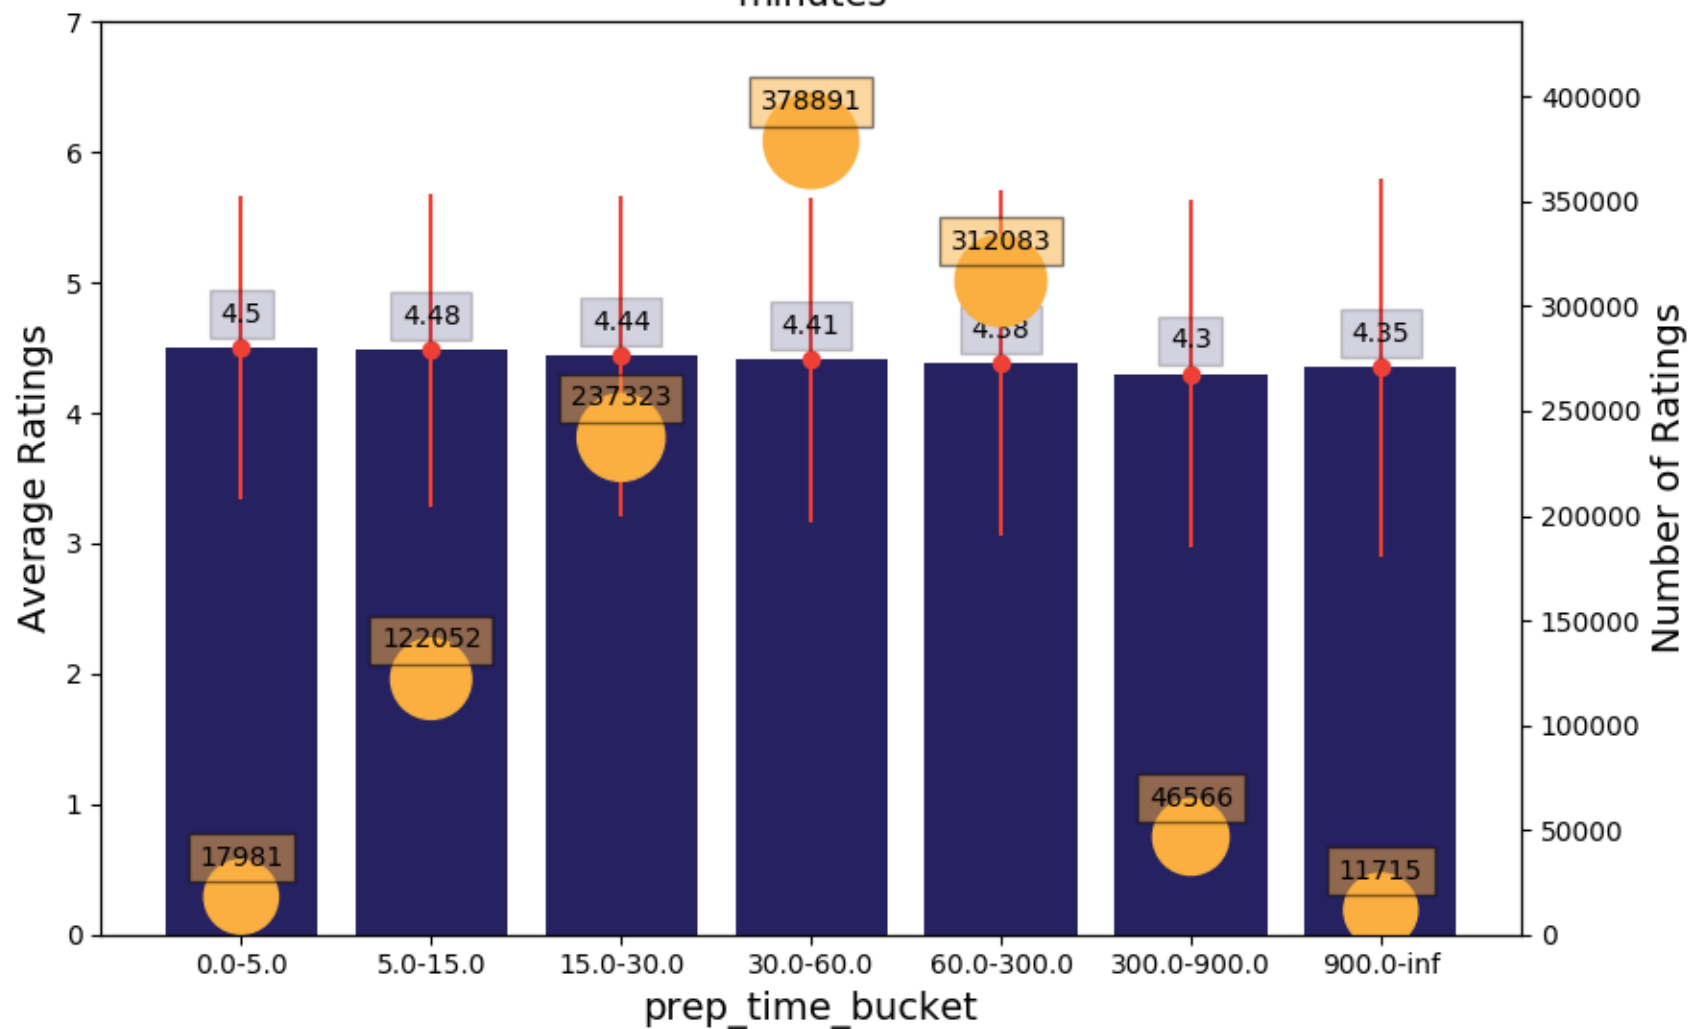
user_id	recipe_id	review_date	submitted	days_since_submission_on_review_date	months_since_submission_on_review_date	years_since_submission_on_review_date
38094	40893	2003-02-17 00:00:00	2002-09-21	149	4.87096774	0.40591397833333337
1293707	40893	2011-12-21 00:00:00	2002-09-21	3378	111.0	9.25
8937	44394	2002-12-01 00:00:00	2002-10-27	35	1.16129032	0.09677419333333333
1982632	54638	2011-08-23 00:00:00	2003-02-23	3103	102.0	8.5
627232	44239	2008-01-20 00:00:00	2002-10-25	1913	62.83870968	5.23655914

only showing top 5 rows

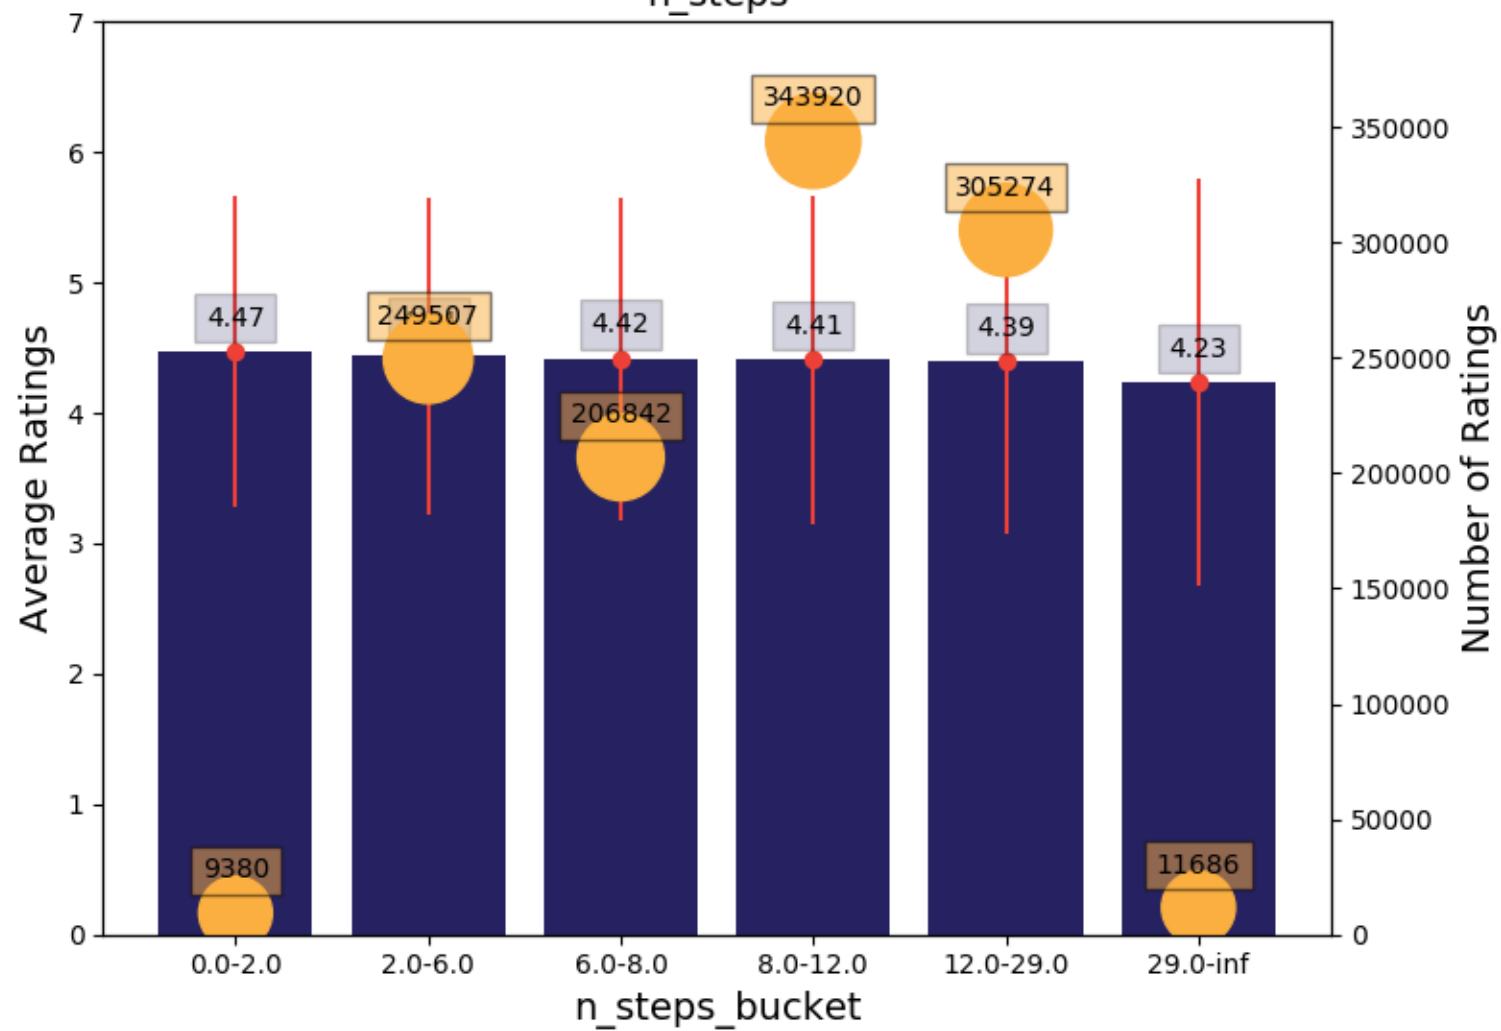
Task 7: Processing Numerical Columns & Exploratory Data Analysis



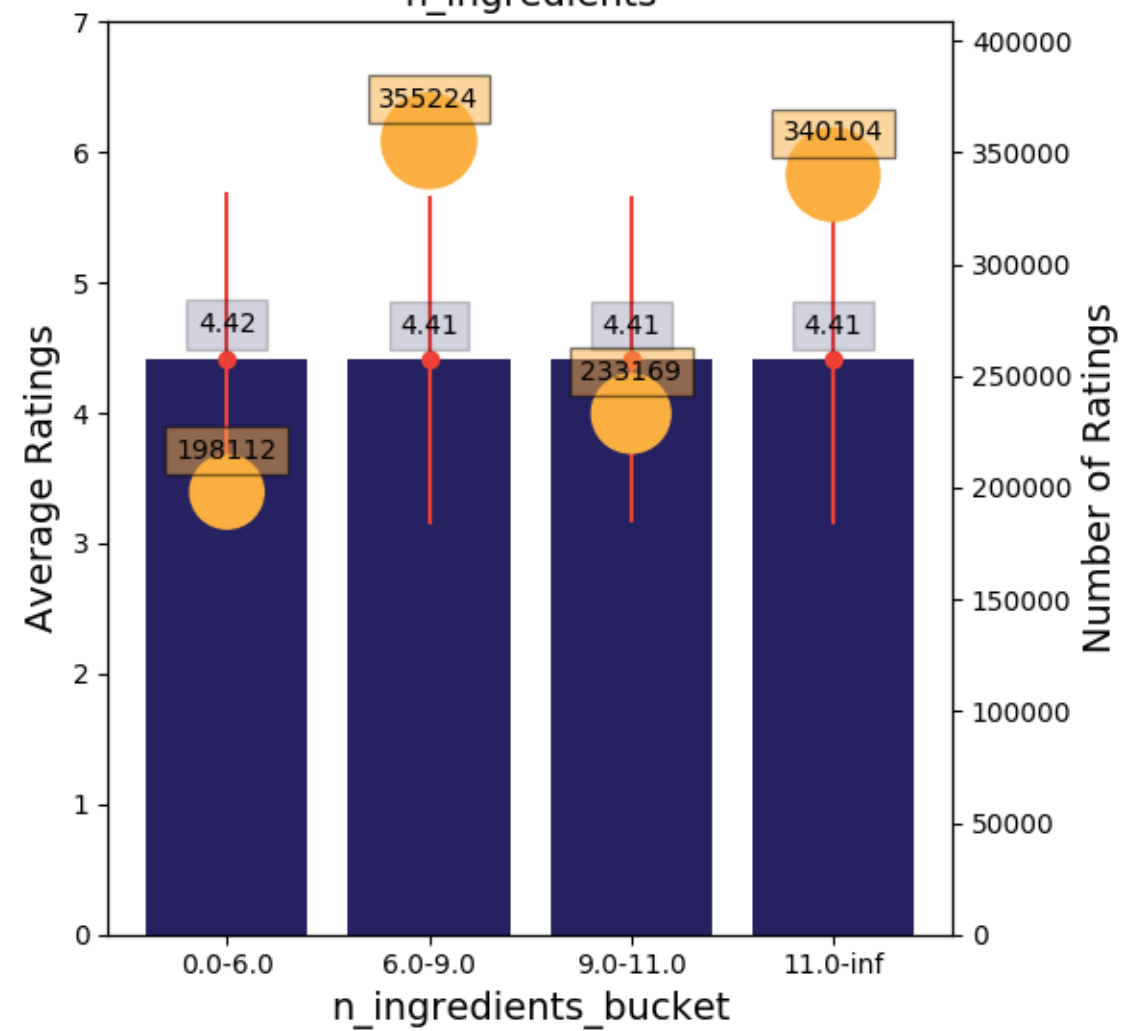
Bucketwise average ratings and number of ratings for minutes



Bucketwise average ratings and number of ratings for
n_steps



Bucketwise average ratings and number of ratings for
n_ingredients




```
In [41]: nutrition_col_quantile_summary
```

	calories	total_fat_PDV	sugar_PDV	sodium_PDV	protein_PDV	\
0.00-0.25	4.416167	4.393560	4.416368	4.423843	4.422679	
0.25-0.50	4.428239	4.420946	4.434610	4.408631	4.419234	
0.50-0.75	4.418471	4.427758	4.403005	4.422298	4.410517	
0.75-0.95	4.393681	4.411867	4.406985	4.397199	4.399930	
0.95 - 1.00	4.342026	4.371152	4.332979	4.373164	4.372771	

	saturated_fat_PDV	carbohydrates_PDV	total_fat_per_100_cal	\
0.00-0.25	4.396849	4.439453	4.385130	
0.25-0.50	4.422079	4.421594	4.396544	
0.50-0.75	4.423550	4.417625	4.415027	
0.75-0.95	4.414637	4.382006	4.438629	
0.95 - 1.00	4.351969	4.324922	4.476848	

	sugar_per_100_cal	sodium_per_100_cal	protein_per_100_cal	\
0.00-0.25	4.412924	4.417770	4.412914	
0.25-0.50	4.427622	4.398280	4.413150	
0.50-0.75	4.410983	4.423546	4.416714	
0.75-0.95	4.392021	4.414099	4.402535	
0.95 - 1.00	4.413383	4.389711	4.404998	

	saturated_fat_per_100_cal	carbohydrates_per_100_cal
0.00-0.25	4.394655	4.438210
0.25-0.50	4.405957	4.419843
0.50-0.75	4.412177	4.400808
0.75-0.95	4.429877	4.379254
0.95 - 1.00	4.446134	4.399911

Activate Windows
Go to Settings to activate

Columns to be bucketized.

1. years_since_submission_on_review_date
2. minutes
3. calories
4. total_fat_PDV
5. sugar_PDV
6. sodium_PDV
7. protein_PDV
8. saturated_fat_PDV
9. carbohydrates_PDV

After creating buckets, study the variation of the average rating for each bucket and decide whether or not a particular bucketed column should be kept in the analysis.

Task 8: Create user-level features

create the following user-level features:

- user_avg_rating
- user_avg_n_ratings
- user_avg_years_betwn_review_and_submission
- user_avg_prep_time_recipes_reviewed
- user_avg_n_steps_recipes_reviewed
- user_avg_n_ingredients_recipes_reviewed
- user_avg_years_betwn_review_and_submission_high_ratings
- user_avg_calories_recipes_reviewed
- user_avg_total_fat_per_100_cal_recipes_reviewed
- user_avg_sugar_per_100_cal_recipes_reviewed
- user_avg_sodium_per_100_cal_recipes_reviewed
- user_avg_protein_per_100_cal_recipes_reviewed
- user_avg_saturated_fat_per_100_cal_recipes_reviewed
- user_avg_carbohydrates_per_100_cal_recipes_reviewed
- user_avg_prep_time_recipes_reviewed_high_ratings
- user_avg_n_steps_recipes_reviewed_high_ratings
- user_avg_n_ingredients_recipes_reviewed_high_ratings

Here, high ratings refer to only those reviews where the user has given five ratings to a recipe.

Adding user level average features

```
In [1]: 1 partition = Window.partitionBy("user_id")
2
3 interaction_level_df = (interaction_level_df
4     .withColumn("user_avg_rating",
5         F.avg(F.col("rating")).over(partition))
6     .withColumn("user_n_ratings",
7         F.count(F.col("rating")).over(partition))
8     .withColumn("user_avg_years_betwn_review_and_submission",
9         F.avg(F.col("years_since_submission_on_review_date")).over(partition))
10    .withColumn("user_avg_prep_time_recipes_reviewed",
11        F.avg(F.col("minutes")).over(partition))
12    .withColumn("user_avg_n_steps_recipes_reviewed",
13        F.avg(F.col("n_steps")).over(partition))
14    .withColumn("user_avg_n_ingredients_recipes_reviewed",
15        F.avg(F.col("n_ingredients")).over(partition)))
```

More Features:

high_ratings = 5 rating

- user_avg_years_betwn_review_and_submission_high_ratings
- user_avg_prep_time_recipes_reviewed_high_ratings
- user_avg_n_steps_recipes_reviewed_high_ratings
- user_avg_n_ingredients_recipes_reviewed_high_ratings

```
In [18]: interaction_level_df = (interaction_level_df
1     .withColumn("ind_5_rating",
2         F.when(interaction_level_df["rating"] != 5, None)
3         .otherwise(1))
4     .withColumn("years_since_submission_on_review_date_5_ratings",
5         F.when(interaction_level_df["rating"] != 5, None)
6         .otherwise(F.col("years_since_submission_on_review_date")))
7     .withColumn("minutes_5_ratings",
8         F.when(interaction_level_df["rating"] != 5, None)
9         .otherwise(F.col("minutes")))
10    .withColumn("n_steps_5_ratings",
11        F.when(interaction_level_df["rating"] != 5, None)
12        .otherwise(F.col("n_steps")))
13    .withColumn("n_ingredients_5_ratings",
14        F.when(interaction_level_df["rating"] != 5, None)
15        .otherwise(F.col("n_ingredients")))
```

Activate Windows
Go to Settings to activate Windows

Task 9: Create tag-level features

Extract tags-level features. If you extract and list unique tags and explore all the available tags, you will realize that tags hold a lot of information about the recipe.

For example, the healthy tag signifies that the person who uploaded the recipe considers it healthy.

If a user specifically looks for the healthy tag, you would want to recommend more healthy recipes to them.

Find the most value-adding tags and create features to capture them.

1. Top `n` most rated tags

```
In [27]: tags_ratings_summary.sort(F.col("n_user_ratings").desc()).show(20)
```

► Spark Job Progress

individual_tag	avg_user_rating	n_user_ratings	n_recipes	in_percent_recipes	in_percent_interactions
preparation	4.411751277206117	1121393	228634	0.9923092280582971	0.9953701772309648
course	4.414709994170975	1055065	212023	0.9202147513519613	0.9364961579394449
time-to-make	4.42448745887648	927389	183484	0.7963507894759685	0.8231684639480068
dietary	4.411783400011269	887350	160286	0.6956676475439008	0.7876290709554069
main-ingredient	4.424306327204302	863051	169236	0.7345121220801542	0.7660608072543358
easy	4.418303138271644	628690	125028	0.5426421186948257	0.5580374380108805
occasion	4.414476975563467	619646	113426	0.49228752723453384	0.5500098081943248
cuisine	4.416987941239125	478822	90622	0.3933144102150118	0.42501169438554104
low-in-something	4.41607825652905	412694	76654	0.33269098894994054	0.36631519897320186
main-dish	4.3960733455628125	383227	71230	0.30914993533154517	0.34015971823409896
equipment	4.42395792662005	338076	48336	0.20978620348428426	0.30008281488963784
[60-minutes-or-less	4.405319536361468	318524	64042	0.27795283108946817	0.282728080460923
meat	4.408245836621744	300297	50669	0.21991180785222608	0.26654944173178097
taste-mood	4.412394148815225	290266	47262	0.20512486653993386	0.25764573157146803
north-american	4.413212293557913	283433	48182	0.20911781811237554	0.25158062823925603
vegetables	4.45447178628346	259147	53344	0.23152174856557556	0.2300239035903317
oven	4.417805174050443	249669	30777	0.1335772505924325	0.22161104695595366
[30-minutes-or-less	4.424764743868313	246221	50347	0.21851427480187147	0.2185505352788767
4-hours-or-less]	4.385306381697524	242459	48028	0.20844943274046682	0.2152113111114859

Activate Windows
Go to Settings to activate

2. Bottom n least rated tags

```
In [34]: tags_ratings_summary.sort(F.col("n_user_ratings").asc()).show(5)
```

► Spark Job Progress

individual_tag	avg_user_rating	n_user_ratings	n_recipes	in_percent_recipies	in_percent_interactions
spaghetti-sauce]	5.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]
spaghetti-sauce]	4.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]
main-dish-seafood]	0.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]
roast-beef-comfor...	5.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]
beans-side-dishes]	5.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]

only showing top 5 rows

The above tags are present in 1 recipe in over two hundred thousand. The features we create based on these tags will not teach the model new information. If these tags were one hot encoded, the entire column would be filled with zeros, and only a few rows will have 1s. One hot encoding of these tags is not a good idea. If you come up with an encoding that captures the rarity of these tags, only then can you add these tags to the analysis.

3. Top n rated tags

```
In [35]: tags_ratings_summary.sort(F.col("avg_user_rating").desc()).show(5)
```

► Spark Job Progress

individual_tag	avg_user_rating	n_user_ratings	n_recipes	in_percent_recipies	in_percent_interactions
side-dishes-beans]	5.0]	2]	2]	8.680329505308021E-6]	1.775238791807983E-6]
[healthy]	5.0]	4]	3]	1.302049425796203...	3.550477583615966E-6]
cranberry-sauce]	5.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]
breakfast-potatoes]	5.0]	1]	1]	4.340164752654011E-6]	8.876193959039915E-7]
occasion]	5.0]	3]	1]	4.340164752654011E-6]	2.662858187711975E-6]

only showing top 5 rows

Top rated tags have low number of ratings.